



V8中的trampoline和Embedded builtins

陆亚涵

中科院软件所 PLCT实验室

2021-05-28

1.Trampoline机制

2.Embedded builtins

3.Short builtin calls

1.Trampoline机制

1.Trampoline机制

Backgroud

riscv64 example:

```
1      addi  a0, a1, 1
2
3      bne   a0, zero_reg, 4
4
5      j     L1
6
7  L2:  addi  a0, a1, 1
8
9      addi  a0, a2, 2
10
11     ...
12
13     j     L2
14
15     ...
16  L1 : ret
```

1.Trampoline机制

Backgroud

riscv64 example:

```
1      addi  a0, a1, 1
2
3      bne   a0, zero_reg, 4
4
5      j     L1
6
7 L2:    addi  a0, a1, 1
8
9      addi  a0, a2, 2
10
11     ...
12
13     j     L2
14
15     ...
16 L1 : ret
```

对于 Gcc/LLVM等编译静态语言如c++/c, 可以扫描整个汇编程序后确定j的offset后再选择将伪指令j汇编成near jump 还是far jump

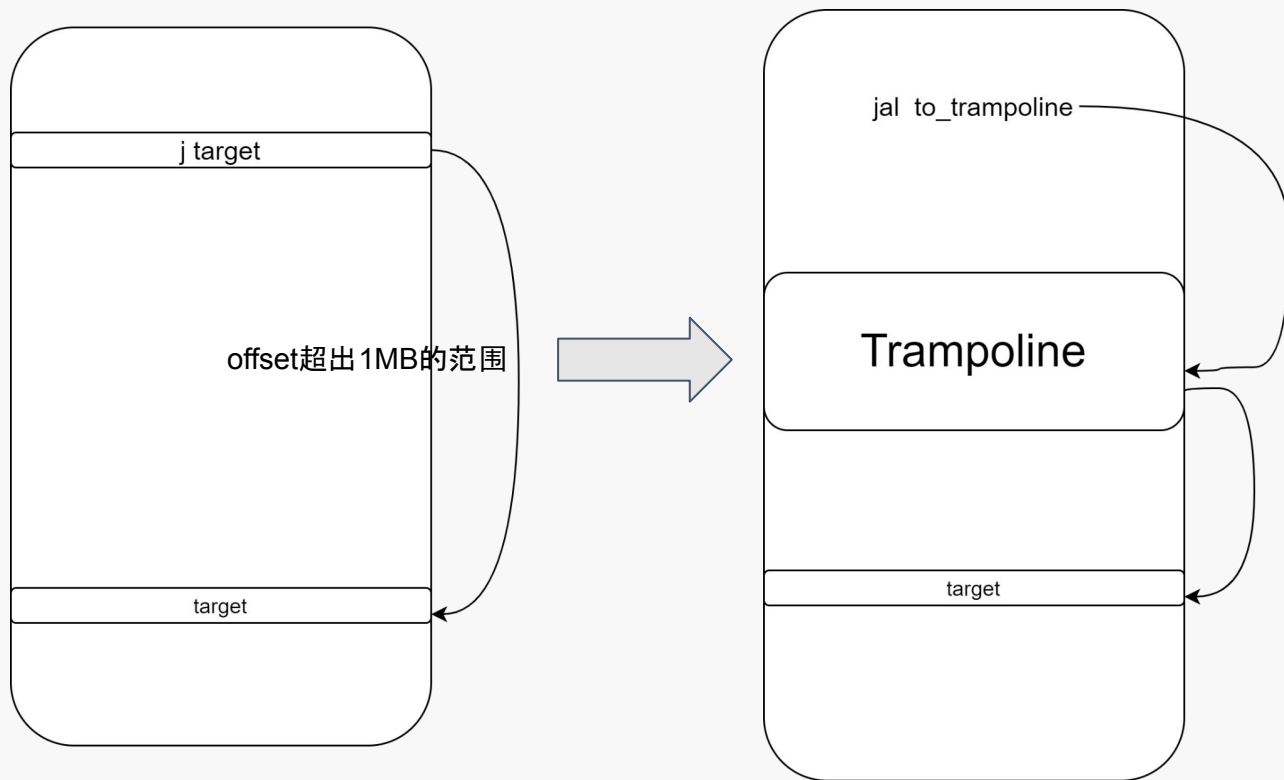
但对于JS语言来说,V8等引擎都是对JS代码进行JIT编译的, 只会对V8生成的IR扫描一次就生成 汇编代码.

表1 不同架构直接跳转的范围

x86	arm64	mips64r6	riscv64	mips64r2
Jmp: $\pm 2\text{GB}$	B/BL: $\pm 128\text{MB}$	Bc: $\pm 128\text{MB}$	jal:($\pm 1\text{ MiB}$)	B: $\pm 128\text{KB}$
	B.EQ: $\pm 1\text{MB.}$	BEQ:($\pm 32\text{KB}$)	beq:($\pm 4096\text{ B}$)	BEQ:($\pm 32\text{KB}$)

1.Trampoline机制

Backgroud



1.Trampoline机制

V8实现

```
void TurboAssembler::Branch(Label* L, Condition cond, Register rs,
                             const Operand& rt) {
    if (L->is_bound()) {
        if (!BranchShortCheck(0, L, cond, rs, rt)) {
            if (cond != cc_always) {
                Label skip;
                Condition neg_cond = NegateCondition(cond);
                BranchShort(&skip, neg_cond, rs, rt);
                BranchLong(L);
                bind(&skip);
            } else {
                BranchLong(L);
                EmitConstPoolWithJumpIfNeeded();
            }
        }
    } else {
        if (is_trampoline_emitted()) {
            if (cond != cc_always) {
                Label skip;
                Condition neg_cond = NegateCondition(cond);
                BranchShort(&skip, neg_cond, rs, rt);
                BranchLong(L);
                bind(&skip);
            } else {
                BranchLong(L);
                EmitConstPoolWithJumpIfNeeded();
            }
        } else {
            BranchShort(L, cond, rs, rt);
        }
    }
}
```

1.Trampoline机制

实现

```
void TurboAssembler::BranchShortHelper(int32_t offset, Label* L) {
    DCHECK(L == nullptr || offset == 0);
    offset = GetOffset(offset, L, OffsetSize::kOffset21);
    j(offset);
}

void TurboAssembler::BranchShort(int32_t offset) {
    DCHECK(is_int21(offset));
    BranchShortHelper(offset, nullptr);
}

void TurboAssembler::BranchShort(Label* L) { BranchShortHelper(0, L); }

int32_t TurboAssembler::GetOffset(int32_t offset, Label* L, OffsetSize bits) {
    if (L) {
        offset = branch_offset_helper(L, bits);
    } else {
        DCHECK(is_intn(offset, bits));
    }
    return offset;
}
```

```
int32_t Assembler::branch_offset_helper(Label* L, OffsetSize bits) {
    int32_t target_pos;

    DEBUG_PRINTF("branch_offset_helper: %p to %p (%d)\n", L,
                 reinterpret_cast<Instr*>(buffer_start_ + pc_offset()),
                 pc_offset());

    if (L->is_bound()) {
        target_pos = L->pos();
        DEBUG_PRINTF("\tbound: %d", target_pos);
    } else {
        if (L->is_linked()) {
            target_pos = L->pos();
            L->link_to(pc_offset());
            DEBUG_PRINTF("\tadded to link: %d\n", target_pos);
        } else {
            L->link_to(pc_offset());
            if (!trampoline_emitted_) {
                unbound_labels_count++;
                next_buffer_check_ -= kTrampolineSlotsSize;
            }
            DEBUG_PRINTF("\tstarted link\n");
            return kEndOfJumpChain;
        }
    }

    int32_t offset = target_pos - pc_offset();
    DCHECK(is_intn(offset, bits));
    DCHECK_EQ(offset & 1, 0);
    DEBUG_PRINTF("\toffset = %d\n", offset);
    return offset;
}
```


1.Trampoline机制

V8实现

```
void Assembler::emit(Instr x) {  
    if (!is_buffer_growth_blocked()) {  
        CheckBuffer();  
    }  
    DEBUG_PRINTF("%p: ", pc_);  
    disassembleInstr(x);  
    EmitHelper(x);  
    CheckTrampolinePoolQuick();  
}
```

```
void CheckTrampolinePoolQuick(int extra_instructions = 0) {  
    DEBUG_PRINTF("\tpc_offset:%d %d\n", pc_offset(),  
        next_buffer_check_ - extra_instructions * kInstrSize);  
    if (pc_offset() >= next_buffer_check_ - extra_instructions * kInstrSize) {  
        CheckTrampolinePool();  
    }  
}
```

```
// For BROCKTrampolinePoolScope buffer.
```

```
next_buffer_check_ = FLAG_force_long_branches
```

```
    ? kMaxInt
```

```
    : kMaxBranchOffset - kTrampolineSlotsSize * 16;
```

实现

```

if (unbound_labels_count_ > 0) {
    // First we emit jump, then we emit trampoline pool.
    {
        DEBUG_PRINTF("inserting trampoline pool at %p (%d)\n",
            reinterpret_cast<Instr*>(buffer_start_ + pc_offset()),
            pc_offset());
        BlockTrampolinePoolScope block_trampoline_pool(this);
        Label after_pool;
        j(&after_pool);

        int pool_start = pc_offset();
        for (int i = 0; i < unbound_labels_count_; i++) {
            int64_t imm64;
            imm64 = branch_long_offset(&after_pool);
            DCHECK(is_int32(imm64));
            int32_t Hi20 = (((int32_t)imm64 + 0x800) >> 12);
            int32_t Lo12 = (int32_t)imm64 << 20 >> 20;
            auipc(t6, Hi20); // Read PC + Hi20 into t6
            jr(t6, Lo12); // jump PC + Hi20 + Lo12
        }
        // If unbound_labels_count_ is big enough, label after_pool will
        // need a trampoline too, so we must create the trampoline before
        // the bind operation to make sure function 'bind' can get this
        // information.
        trampoline_ = Trampoline(pool_start, unbound_labels_count_);
        bind(&after_pool);

        trampoline_emitted_ = true;
        // As we are only going to emit trampoline once, we need to prevent any
        // further emission.
        next_buffer_check_ = kMaxInt;
    }
} else {
    // Number of branches to unbound label at this point is zero, so we can
    // move next buffer check to maximum.
    next_buffer_check_ =
        pc_offset() + kMaxBranchOffset - kTrampolineSlotsSize * 16;
}

return;

```

1.Trampoline机制


实现

-- B17 start --				
0x444328246c	4ec	02071293	slli	t0, a4, 32
0x4443282470	4f0	7f8b0d13	addi	s10, s6, 2040
0x4443282474	4f4	768d3383	ld	t2, 1896(s10)
0x4443282478	4f8	02039393	slli	t2, t2, 32
0x444328247c	4fc	227284e3	beq	t0, t2, 2600 → 0x4443282ea4 (Call_ReceiverIsNullOrUndefined_Baseline_Compact)
B18 start				
0x4443282ea0	f20	0740006f	j	116 → 0x4443282f14 (Call_ReceiverIsNullOrUndefined_Baseline_Compact)
0x4443282ea4	f24	00000f97	auipc	t6, 0x0
0x4443282ea8	f28	698f8067	jalr	zero_reg, 1688(t6) → 0x444328353c (Call_ReceiverIsNullOrUndefined_Baseline_Compact)
0x4443282eac	f2c	00001f97	auipc	t6, 0x1
0x4443282eb0	f30	2c8f8067	jalr	zero_reg, 712(t6) → 0x4443284174 (Call_ReceiverIsNullOrUndefined_Baseline_Compact)
0x4443282eb4	f34	00001f97	auipc	t6, 0x1
0x4443282eb8	f38	334f8067	jalr	zero_reg, 820(t6) → 0x44432841e8 (Call_ReceiverIsNullOrUndefined_Baseline_Compact)
0x4443282ebc	f3c	00001f97	auipc	t6, 0x1
0x4443282ec0	f40	330f8067	jalr	zero_reg, 816(t6) → 0x44432841ec (Call_ReceiverIsNullOrUndefined_Baseline_Compact)
0x4443282ec4	f44	00001f97	auipc	t6, 0x1
0x4443282ec8	f48	354f8067	jalr	zero_reg, 852(t6) → 0x4443284218 (Call_ReceiverIsNullOrUndefined_Baseline_Compact)
0x4443282ecc	f4c	00001f97	auipc	t6, 0x1
0x4443282ed0	f50	378f8067	jalr	zero_reg, 888(t6) → 0x4443284244 (Call_ReceiverIsNullOrUndefined_Baseline_Compact)
0x4443282ed4	f54	00001f97	auipc	t6, 0x1
0x4443282ed8	f58	428f8067	jalr	zero_reg, 1064(t6) → 0x44432842fc (Call_ReceiverIsNullOrUndefined_Baseline_Compact)
0x4443282edc	f5c	00001f97	auipc	t6, 0x1
0x4443282ee0	f60	44cf8067	jalr	zero_reg, 1100(t6) → 0x4443284328 (Call_ReceiverIsNullOrUndefined_Baseline_Compact)
0x4443282ee4	f64	00001f97	auipc	t6, 0x1
0x4443282ee8	f68	470f8067	jalr	zero_reg, 1136(t6) → 0x4443284354 (Call_ReceiverIsNullOrUndefined_Baseline_Compact)
0x4443282eec	f6c	00001f97	auipc	t6, 0x1
0x4443282ef0	f70	54cf8067	jalr	zero_reg, 1356(t6) → 0x4443284438 (Call_ReceiverIsNullOrUndefined_Baseline_Compact)
0x4443282ef4	f74	00000f97	auipc	t6, 0x0
0x4443282ef8	f78	020f8067	jalr	zero_reg, 32(t6) → 0x4443282f14 (Call_ReceiverIsNullOrUndefined_Baseline_Compact)
-- B94 start --				
0x444328353c	15bc	fe843b83	ld	s7, -24(fp)
[DecompressTaggedPointer				
0x4443283540	15c0	ffffbe703	lwu	a4, -1(s7)
0x4443283544	15c4	00ed8733	add	a4, s11, a4
]				
0x4443283548	15c8	00010993	mv	s3, sp
0x444328354c	15cc	ff810113	addi	sp, sp, -8
0x4443283550	15d0	ff017113	andi	sp, sp, 0xfffffffff0
0x4443283554	15d4	01313023	sd	s3, 0(sp)
0x4443283558	15d8	fce43023	sd	a4, -64(fp)

2. Embedded builtins

2.Embedded builtins

Background

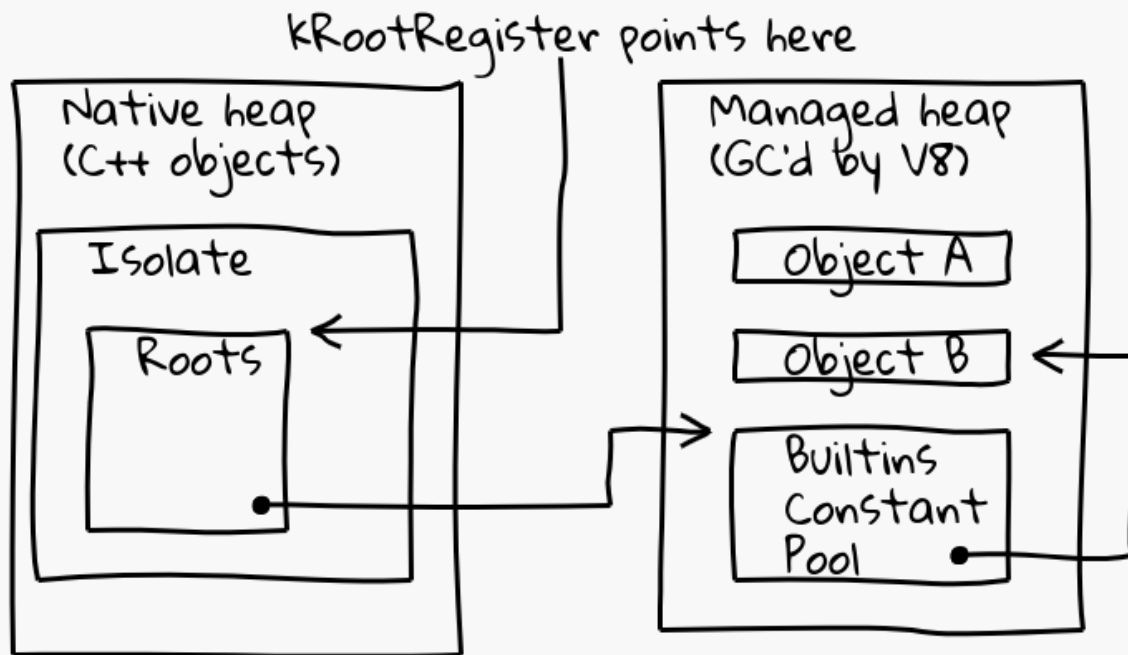


V8 snapshot size (including builtins) from 2015 until 2017

<https://v8.dev/blog/embedded-builtins>

2018年,spectre漏洞的发现, Chrome 浏览器开启了 site isolation(站点隔离)特性. 因此造成浏览器中的每一个网页都会创建大量的渲染线程和更多的V8 isolates实例,每个实例都包含一份 builtin代码.

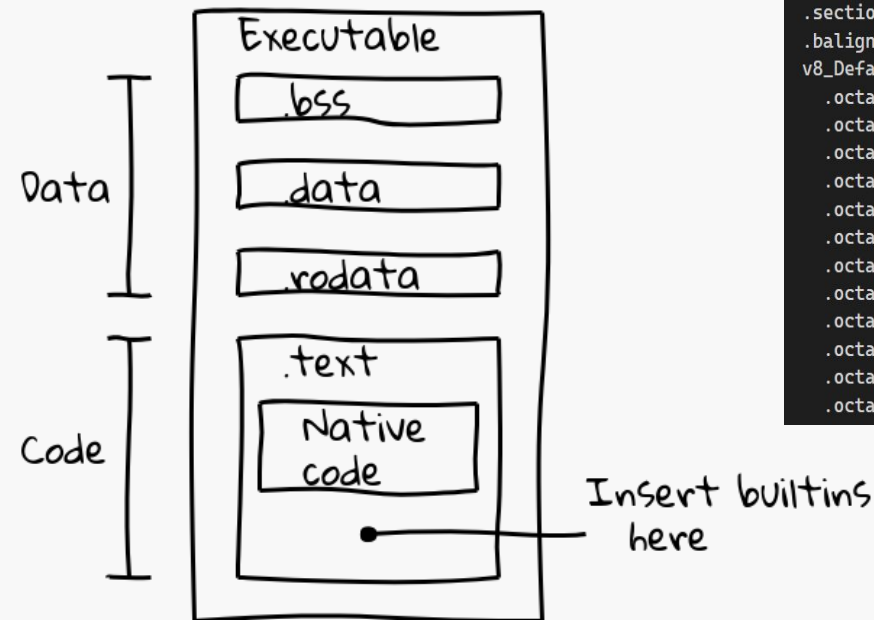
2.Embedded builtins



Isolate- and process-independent code

<https://v8.dev/blog/embedded-builtins>

2.Embedded builtins



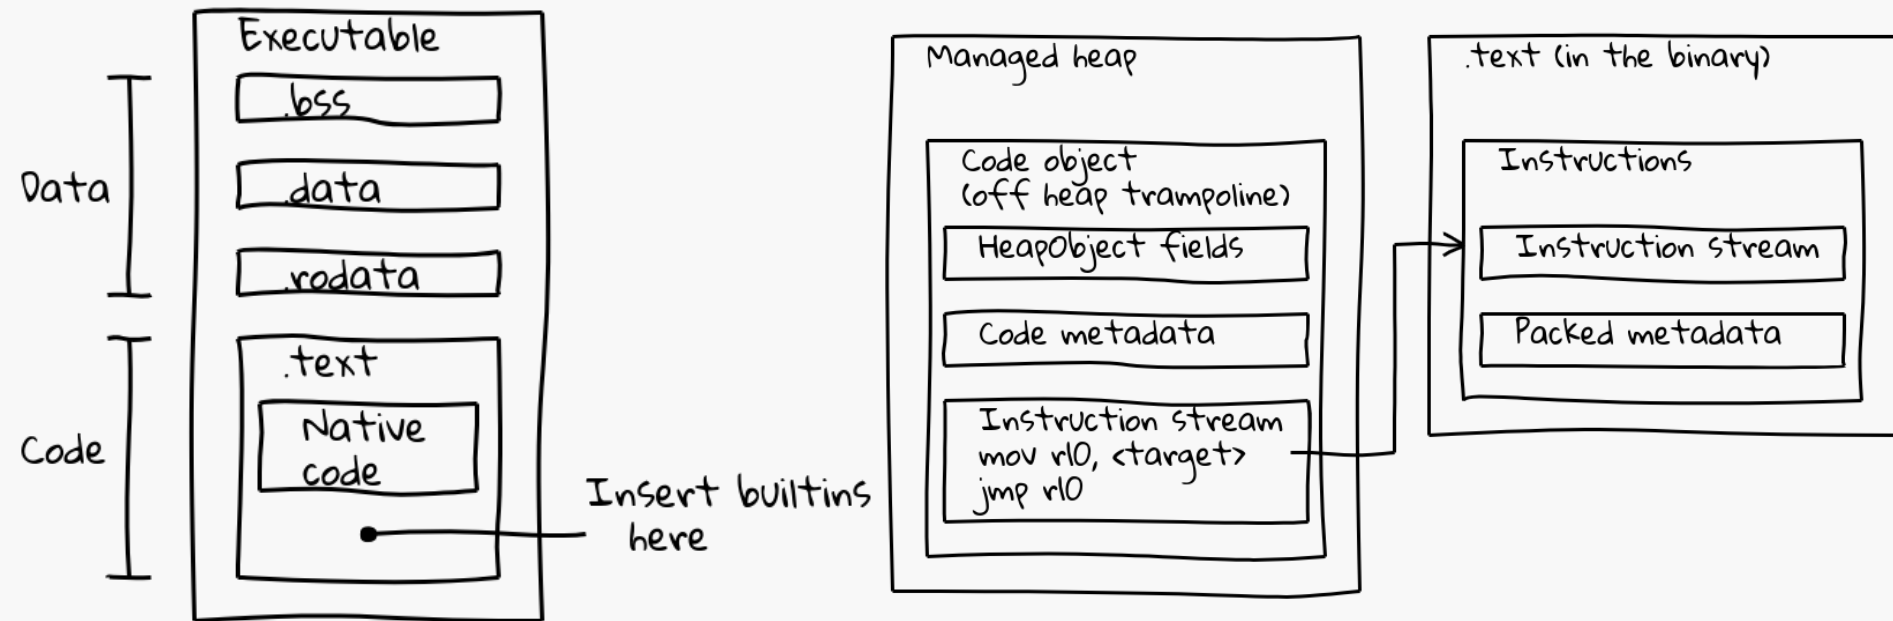
```
// The embedded blob data section starts here.  
.section .rodata  
.balign 8  
v8_Default_embedded_blob_data_data_  
  .octa 0xc790027f000000003bcaf050,0x5a0000000003ae53c50d042f227,0x218000005c00000037800000000  
  .octa 0x3a8000007e00000016c00000378,0x17000000ba000000378000004e4  
  .octa 0x19000000d200000016c0000085c,0xe800000ec000000088000009c8  
  .octa 0x26000000fc0000008800000a50,0x2ac000012400000031500000ad8  
  .octa 0x324000015000000033c00000df0,0x35000001840000003430000112c  
  .octa 0x12000001ba0000003c700001470,0xc000001ce0000022400001838  
  .octa 0xc000001dc00000011600001a5c,0xc000001ea00000011600001b74,0x2b7c00001f80000011600001c8c  
  .octa 0x2b6c00004b0000003e4a00001da4,0x2b8c0000768000003e4a00005bf0  
  .octa 0x2b7c0000a22000003e4a00009a3c,0x2b8c0000cda000003e4a0000d888  
  .octa 0x2b7c0000f94000003e4a000116d4,0x2bb0000124c000003e4a00015520  
  .octa 0x2bb00001508000003e760001936c,0x2bb000017c4000003e760001d1e4  
  .octa 0xfd00001a80000003e760002105c,0xd80001b7e000001a8a00024ed4
```

V8 mksnapshot生成的emdedd.S文件

Sections of an executable binary file

<https://v8.dev/blog/embedded-builtins>

2.Embedded builtins

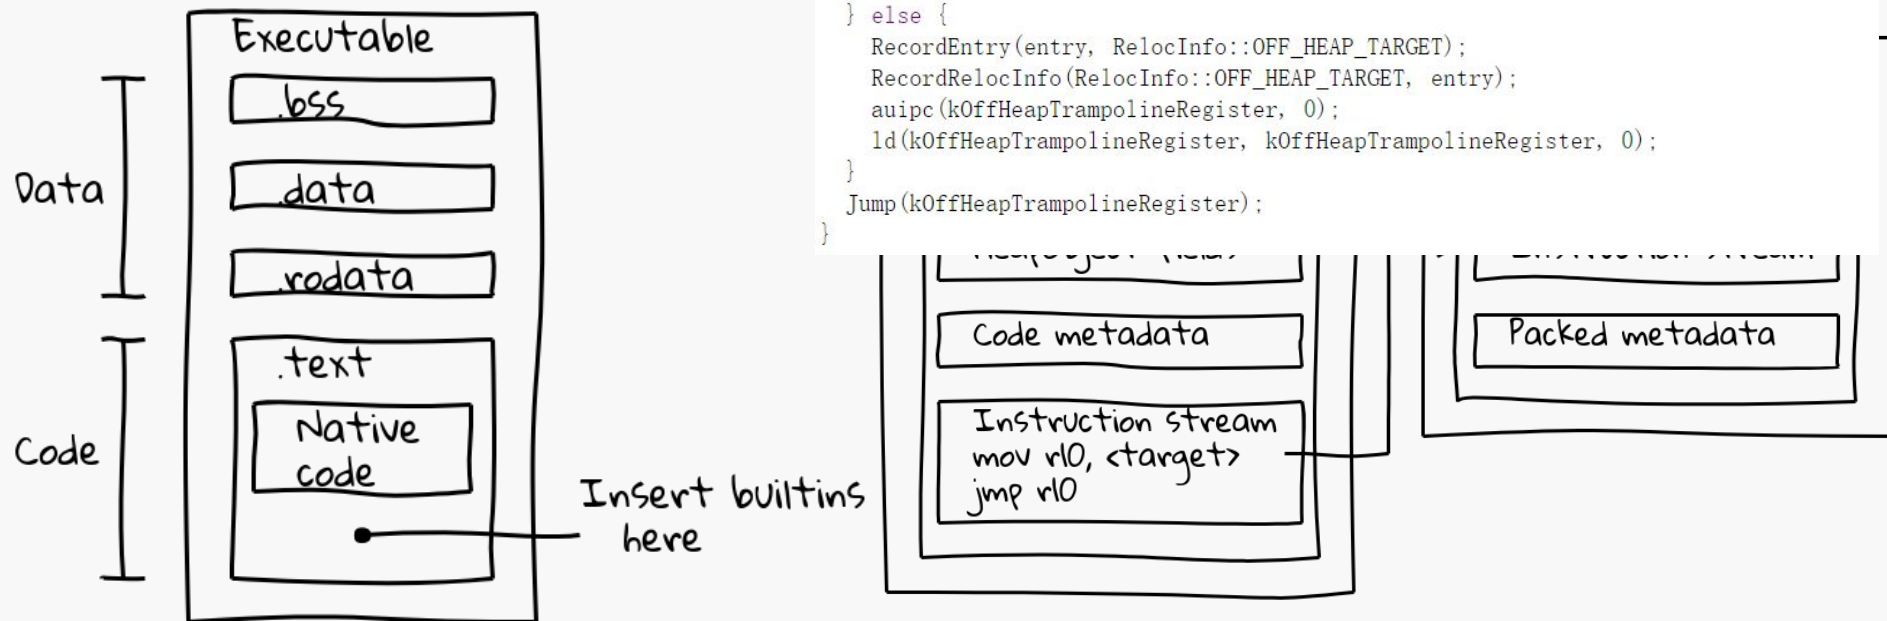


Sections of an executable binary file

<https://v8.dev/blog/embedded-builtins>

2.Embedded builtins

```
void MacroAssembler::JumpToInstructionStream(Address entry) {  
    // Ld a Address from a constant pool.  
    // Record a value into constant pool.  
    if (!FLAG_riscv_constant_pool) {  
        li(kOffHeapTrampolineRegister, Operand(entry, RelocInfo::OFF_HEAP_TARGET));  
    } else {  
        RecordEntry(entry, RelocInfo::OFF_HEAP_TARGET);  
        RecordRelocInfo(RelocInfo::OFF_HEAP_TARGET, entry);  
        auipc(kOffHeapTrampolineRegister, 0);  
        ld(kOffHeapTrampolineRegister, kOffHeapTrampolineRegister, 0);  
    }  
    Jump(kOffHeapTrampolineRegister);  
}
```



Sections of an executable binary file

<https://v8.dev/blog/embedded-builtins>

3.Short builtin calls

3.Short builtin calls

Due to Spectre v2 various device/OS combinations have turned off indirect branch prediction. This means that on such configurations we'll get very costly stalls on function calls from JIT code that rely on the CallFunction builtin

Spectre : variant NO.2

- exploit Indirect Branches misprediction

victim:
pc with virtual address A: jump good-register (or call good-register)
attacker:
pc with virtual address equal or alias to A: jump gadget-register(or call gadget-register)
gadget code:
transiently read security info and code it into side-channel

Fig. 2: The branch predictor is (mis-)trained in the attacker-controlled context A. In context B, the branch predictor makes its prediction on the basis of training data from context A, leading to speculative execution at an attacker-chosen address which corresponds to the location of the Spectre gadget in the victim's address space.

-- "Spectre Attacks: Exploiting Speculative Execution"

Spectre v2 攻击说明

bilibili:BV1hp4y1t7Mx

For 64-bit applications, branch prediction performance can be negatively impacted when the target of a branch is more than 4 GB away from the branch.